

NAME: JITESH VINOD ZOPE

ROLL NO.: 69

PID: 246054

BATCH: 3

TITLE: IMPLEMENTATION OF AN END-TO-END MACHINE LEARNING DATA PIPELINE

```
In [58]: # pip install numpy
# pip install pandas
# pip install matplotlib
!pip install numpy scikit-learn
!pip install numpy seaborn
```

```
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: numpy in c:\users\hp\appdata\roaming\python\python39
\site-packages (2.0.2)
Requirement already satisfied: scikit-learn in c:\users\hp\appdata\roaming\python\py
thon39\site-packages (1.6.1)
Requirement already satisfied: joblib>=1.2.0 in c:\users\hp\appdata\roaming\python\p
ython39\site-packages (from scikit-learn) (1.5.3)
Requirement already satisfied: scipy>=1.6.0 in c:\users\hp\appdata\roaming\python\py
thon39\site-packages (from scikit-learn) (1.13.1)
Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\hp\appdata\roaming\p
ython\python39\site-packages (from scikit-learn) (3.6.0)
```

```
WARNING: You are using pip version 20.2.3; however, version 25.3 is available.
You should consider upgrading via the 'c:\program files\python39\python.exe -m pip i
nstall --upgrade pip' command.
```

```
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: numpy in c:\users\hp\appdata\roaming\python\python39
\site-packages (2.0.2)
Requirement already satisfied: seaborn in c:\users\hp\appdata\roaming\python\python39\site-packages (0.13.2)
Requirement already satisfied: pandas>=1.2 in c:\users\hp\appdata\roaming\python\python39\site-packages (from seaborn) (2.3.3)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in c:\users\hp\appdata\roaming\python\python39\site-packages (from seaborn) (3.9.4)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\hp\appdata\roaming\python\python39\site-packages (from pandas>=1.2->seaborn) (2.9.0.post0)
Requirement already satisfied: tzdata>=2022.7 in c:\users\hp\appdata\roaming\python\python39\site-packages (from pandas>=1.2->seaborn) (2025.3)
Requirement already satisfied: pytz>=2020.1 in c:\users\hp\appdata\roaming\python\python39\site-packages (from pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: pillow>=8 in c:\users\hp\appdata\roaming\python\python39\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (11.3.0)
Requirement already satisfied: cycler>=0.10 in c:\users\hp\appdata\roaming\python\python39\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\hp\appdata\roaming\python\python39\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.59.1)
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\hp\appdata\roaming\python\python39\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.2.3)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\hp\appdata\roaming\python\python39\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.3.0)
Requirement already satisfied: importlib-resources>=3.2.0; python_version < "3.10" in c:\users\hp\appdata\roaming\python\python39\site-packages (from matplotlib!=3.6.1,
>=3.4->seaborn) (6.5.2)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\hp\appdata\roaming\python\python39\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.7)
Requirement already satisfied: packaging>=20.0 in c:\users\hp\appdata\roaming\python\python39\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (25.0)
Requirement already satisfied: six>=1.5 in c:\users\hp\appdata\roaming\python\python39\site-packages (from python-dateutil>=2.8.2->pandas>=1.2->seaborn) (1.17.0)
Requirement already satisfied: zipp>=3.1.0; python_version < "3.10" in c:\users\hp\appdata\roaming\python\python39\site-packages (from importlib-resources>=3.2.0; python_version < "3.10"->matplotlib!=3.6.1,>=3.4->seaborn) (3.23.0)

WARNING: You are using pip version 20.2.3; however, version 25.3 is available.
You should consider upgrading via the 'c:\program files\python39\python.exe -m pip install --upgrade pip' command.
```

```
In [59]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [60]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
```

```
In [61]: import seaborn as sns
#Load titanic dataset
titanic_data = sns.load_dataset('titanic')
```

```
In [62]: print(titanic_data.shape)
```

```
(891, 15)
```

```
In [63]: print(titanic_data.columns)
```

```
Index(['survived', 'pclass', 'sex', 'age', 'sibsp', 'parch', 'fare',
       'embarked', 'class', 'who', 'adult_male', 'deck', 'embark_town',
       'alive', 'alone'],
      dtype='object')
```

```
In [64]: print(titanic_data.head())
```

```
survived  pclass    sex   age  sibsp  parch     fare embarked  class \
0         0        3  male  22.0     1     0  7.2500      S  Third
1         1        1 female  38.0     1     0 71.2833      C  First
2         1        3 female  26.0     0     0  7.9250      S  Third
3         1        1 female  35.0     1     0 53.1000      S  First
4         0        3  male  35.0     0     0  8.0500      S  Third

      who  adult_male  deck  embark_town  alive  alone
0   man      True    NaN  Southampton   no  False
1 woman     False     C  Cherbourg   yes  False
2 woman     False    NaN  Southampton  yes  True
3 woman     False     C  Southampton  yes  False
4   man      True    NaN  Southampton   no  True
```

```
In [65]: print(titanic_data.tail())
```

```
survived  pclass    sex   age  sibsp  parch     fare embarked  class \
886       0        2  male  27.0     0     0  13.00      S  Second
887       1        1 female  19.0     0     0  30.00      S  First
888       0        3 female  NaN      1     2  23.45      S  Third
889       1        1  male  26.0     0     0  30.00      C  First
890       0        3  male  32.0     0     0   7.75      Q  Third

      who  adult_male  deck  embark_town  alive  alone
886   man      True    NaN  Southampton   no  True
887 woman     False     B  Southampton  yes  True
888 woman     False    NaN  Southampton   no  False
889   man      True     C  Cherbourg   yes  True
890   man      True    NaN  Queenstown  no  True
```

```
In [66]: #Review the structure of the dataset
print(titanic_data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          --          --      
 0   survived    891 non-null   int64  
 1   pclass      891 non-null   int64  
 2   sex         891 non-null   object  
 3   age         714 non-null   float64 
 4   sibsp       891 non-null   int64  
 5   parch       891 non-null   int64  
 6   fare        891 non-null   float64 
 7   embarked    889 non-null   object  
 8   class        891 non-null   category
 9   who          891 non-null   object  
 10  adult_male  891 non-null   bool    
 11  deck         203 non-null   category
 12  embark_town 889 non-null   object  
 13  alive        891 non-null   object  
 14  alone        891 non-null   bool    
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB
None
```

```
In [67]: print(titanic_data.describe())
```

	survived	pclass	age	sibsp	parch	fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
In [68]: #Identifying missing values  
missing_values = titanic_data.isnull().sum()  
print(missing_values)
```

```
survived      0
pclass        0
sex           0
age          177
sibsp        0
parch        0
fare          0
embarked     2
class         0
who           0
adult_male    0
deck          688
embark_town   2
alive         0
alone         0
dtype: int64
```

```
In [69]: #Dealing with missing values  
#Dropping columns with excessive missing data  
new_titanic_df = titanic_data.drop(columns=['deck'])
```

```
In [70]: #Imputing median age for missing age data  
new_titanic_df['age'].fillna(new_titanic_df['age'].median(), inplace=True)
```

C:\Users\HP\AppData\Local\Temp\ipykernel_5248\1960786904.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
new_titanic_df['age'].fillna(new_titanic_df['age'].median(), inplace=True)
```

```
In [71]: missing_values = new_titanic_df.isnull().sum()  
print(missing_values)
```

```
survived      0  
pclass        0  
sex           0  
age           0  
sibsp         0  
parch         0  
fare          0  
embarked      2  
class          0  
who            0  
adult_male     0  
embark_town    2  
alive          0  
alone          0  
dtype: int64
```

```
In [72]: data = new_titanic_df  
data['embark_town'].dtype  
data['embark_town'].unique()  
data['embark_town'].fillna(data['embark_town'].mode()[0], inplace=True)  
data.isnull().sum()
```

```
C:\Users\HP\AppData\Local\Temp\ipykernel_5248\1968151471.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
```

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
data['embark_town'].fillna(data['embark_town'].mode()[0] , inplace=True)
```

```
Out[72]: survived      0
         pclass        0
         sex          0
         age          0
         sibsp        0
         parch        0
         fare          0
         embarked      2
         class         0
         who          0
         adult_male    0
         embark_town   0
         alive         0
         alone         0
         dtype: int64
```

```
In [73]: data = new_titanic_df
data['embarked'].dtype
data['embarked'].unique()
data['embarked'].fillna(data['embarked'].mode()[0] , inplace=True)
data.isnull().sum()
```

```
C:\Users\HP\AppData\Local\Temp\ipykernel_5248\2633234362.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
```

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
data['embarked'].fillna(data['embarked'].mode()[0] , inplace=True)
```

```
Out[73]: survived      0
          pclass        0
          sex           0
          age           0
          sibsp         0
          parch         0
          fare          0
          embarked      0
          class          0
          who            0
          adult_male    0
          embark_town   0
          alive          0
          alone          0
          dtype: int64
```

```
In [74]: # Step 4 : Encode catagorical variables
le = LabelEncoder()
data['sex'] = le.fit_transform(data['sex'])
data['embarked'] = le.fit_transform(data['embarked'])
```

```
In [75]: data.head()
```

```
Out[75]:    survived  pclass  sex  age  sibsp  parch      fare  embarked  class  who  adult_male
0           0       3     1  22.0     1      0    7.2500      2  Third    man    True
1           1       1     0  38.0     1      0   71.2833      0  First   woman   False
2           1       3     0  26.0     0      0    7.9250      2  Third   woman   False
3           1       1     0  35.0     1      0   53.1000      2  First   woman   False
4           0       3     1  35.0     0      0    8.0500      2  Third    man    True
```



```
In [76]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   survived    891 non-null    int64  
 1   pclass      891 non-null    int64  
 2   sex         891 non-null    int64  
 3   age         891 non-null    float64 
 4   sibsp       891 non-null    int64  
 5   parch       891 non-null    int64  
 6   fare         891 non-null    float64 
 7   embarked    891 non-null    int64  
 8   class        891 non-null    category
 9   who          891 non-null    object  
 10  adult_male  891 non-null    bool   
 11  embark_town 891 non-null    object  
 12  alive        891 non-null    object  
 13  alone        891 non-null    bool  
dtypes: bool(2), category(1), float64(2), int64(6), object(3)
memory usage: 79.4+ KB
```

```
In [77]: #Step 5 : Select features and target
data = data[['pclass' , 'sex' , 'age' , 'fare' , 'embarked' , 'survived']]
```

```
X = data[['pclass' , 'sex' , 'age' , 'fare' , 'embarked']]
y = data['survived']
```

```
In [78]: # Step 6 : Train-test split
X_train , X_test,y_train , y_test = train_test_split(X,y,test_size=0.3, random_stat
```

```
In [79]: # Step 7 : Train Model
model = LogisticRegression(max_iter = 1000)
model.fit(X_train,y_train)
```

```
Out[79]: ▾ LogisticRegression ⓘ ?
```

```
LogisticRegression(max_iter=1000)
```

```
In [80]: # Prediction
y_pred = model.predict(X_test)
```

```
In [81]: #from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:",accuracy)
```

Accuracy: 0.7947761194029851

```
In [82]: new_passenger = pd.DataFrame({
    'pclass': [3],
    'sex': ['male'],
    'age' : [28],
    'fare' : [7.25],
    'embarked' : ['S']})
```

```
})
```

```
In [83]: new_passenger_encoded = pd.get_dummies(new_passenger)
new_passenger_encoded = new_passenger_encoded.reindex(columns=X.columns, fill_value=
```

```
In [84]: prediction = model.predict(new_passenger_encoded)
print("Survived" if prediction[0] == 1 else "Did not survive")
```

Survived

```
In [85]: new_passenger = pd.DataFrame({
    'pclass': [1,3,2],
    'sex': ['female','male','female'],
    'age' : [38,45,14],
    'fare' : [80.0, 8.05,20.0],
    'embarked' : ['C','S','Q']

})
```

```
In [86]: new_passenger_encoded = pd.get_dummies(new_passenger)
new_passenger_encoded = new_passenger_encoded.reindex(columns=X.columns, fill_value=
```

```
In [87]: prediction = model.predict(new_passenger_encoded)
print("Survived" if prediction[0] == 1 else "Did not survive")
```

Survived

```
In [88]: predictions = model.predict(new_passenger_encoded)
for i,pred in enumerate(predictions):
    print(f"Passenger {i+1}:", 
          "Survived" if pred == 1 else "Did not survive")
```

Passenger 1: Survived
Passenger 2: Survived
Passenger 3: Survived

TRAINING A MACHINE LEARNING MODEL WITH TITANIC DATASET

```
In [102...]: # Importing necessary libraries
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import seaborn as sns
import pandas as pd

# Loading the Titanic dataset
titanic_df = sns.load_dataset('titanic')

# Splitting the dataset
train_data, test_data = train_test_split(titanic_df, test_size=0.2, random_state=42

print(f"Train data shape: {train_data.shape}")
```

```

print(f"Test data shape: {test_data.shape}")

# Separating target variable
x_train = train_data.drop("survived", axis=1)
y_train = train_data["survived"]

x_test = test_data.drop("survived", axis=1)
y_test = test_data["survived"]

# Convert categorical variables into numerical form
x_train = pd.get_dummies(x_train, drop_first=True)
x_test = pd.get_dummies(x_test, drop_first=True)

# Handling missing values
x_train = x_train.fillna(0)
x_test = x_test.fillna(0)

# Align train and test columns (IMPORTANT)
x_train, x_test = x_train.align(x_test, join="left", axis=1, fill_value=0)

# Initialize Logistic Regression model
logreg = LogisticRegression(max_iter=1000)

# Train the model
logreg.fit(x_train, y_train)

# Make predictions
predictions = logreg.predict(x_test)

# Display evaluation metrics
print("Classification Report:")
print(classification_report(y_test, predictions))

print("Confusion Matrix:")
print(confusion_matrix(y_test, predictions))

print("Accuracy Score:")
print(accuracy_score(y_test, predictions))

```

Train data shape: (712, 15)
Test data shape: (179, 15)
Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	105
1	1.00	1.00	1.00	74
accuracy			1.00	179
macro avg	1.00	1.00	1.00	179
weighted avg	1.00	1.00	1.00	179

Confusion Matrix:

[[105 0]	[0 74]]
----------	----------

Accuracy Score:

1.0

PLOTTING

In [109...]

```
import matplotlib.pyplot as plt
import seaborn as sns

cm = confusion_matrix(y_test, predictions)

plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - Logistic Regression")
plt.show()

comparison_df = pd.DataFrame({
    "Actual": y_test.values,
    "Predicted": predictions
})

comparison_df.value_counts().plot(kind="bar", figsize=(6,4))
plt.title("Actual vs Predicted Survival Counts")
plt.xlabel("(Actual, Predicted)")
plt.ylabel("Count")
plt.show()

plt.figure(figsize=(6,4))
sns.countplot(x=predictions)
plt.xticks([0,1], ["Not Survived", "Survived"])
plt.title("Predicted Survival Distribution")
plt.xlabel("Prediction")
plt.ylabel("Count")
plt.show()

coefficients = pd.DataFrame({
    "Feature": x_train.columns,
    "Coefficient": logreg.coef_[0]
})

coefficients = coefficients.sort_values(by="Coefficient", ascending=False).head(10)

plt.figure(figsize=(8,5))
sns.barplot(x="Coefficient", y="Feature", data=coefficients)
plt.title("Top 10 Important Features (Logistic Regression)")
plt.show()
```


```

from sklearn.metrics import roc_curve, auc


y_prob = logreg.predict_proba(x_test)[:,1]
fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f"AUC = {roc_auc:.2f}")
plt.plot([0,1], [0,1], linestyle="--")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Logistic Regression")
plt.legend()
plt.show()


```





Actual vs Predicted Survival Counts



Predicted Survival Distribution





STANDARDIZATION

In [110]: `from sklearn.preprocessing import StandardScaler`

```

scaler = StandardScaler()

# Standardize the data
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)

logreg = LogisticRegression(max_iter=1000)
logreg.fit(x_train_scaled, y_train)

predictions = logreg.predict(x_test_scaled)

print("Accuracy:", accuracy_score(y_test, predictions))

print("Mean after standardization:", x_train_scaled.mean())
print("Std after standardization:", x_train_scaled.std())

scaler.fit_transform(x_test) # ✗ NEVER fit on test data

```

```

Accuracy: 1.0
Mean after standardization: -6.833810508195883e-18
Std after standardization: 1.0

Out[110... array([[ 0.88742288, -1.34903605,  0.82036305, ..., -0.32394177,
       -1.40830868,  1.19118383],
      [-0.25537349,  0.37905601, -0.55202, ..., -0.32394177,
       0.7100716, -0.83950099],
      [ 0.88742288, -0.23413795, -0.55202, ..., -0.32394177,
       0.7100716, -0.83950099],
      ...,
      [ 0.88742288,  0.76927035,  0.82036305, ..., -0.32394177,
       0.7100716,  1.19118383],
      [-0.25537349, -0.40137266, -0.55202, ..., -0.32394177,
       0.7100716,  1.19118383],
      [ 0.88742288, -1.12605643,  0.82036305, ..., -0.32394177,
       0.7100716,  1.19118383]])

```

NORMALIZATION

```

In [111... from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

# Normalize the data
x_train_norm = scaler.fit_transform(x_train)
x_test_norm = scaler.transform(x_test)

```

```
logreg = LogisticRegression(max_iter=1000)
logreg.fit(x_train_norm, y_train)

predictions = logreg.predict(x_test_norm)

print("Accuracy:", accuracy_score(y_test, predictions))

print("Min value:", x_train_norm.min())
print("Max value:", x_train_norm.max())

scaler.fit_transform(x_test) # ✗ Wrong
```

Accuracy: 1.0

Min value: 0.0

Max value: 1.0

```
Out[111... array([[1.        , 0.        , 0.25      , ..., 0.        , 0.        ,
       1.        ],
      [0.5       , 0.43661972, 0.        , ..., 0.        , 1.        ,
       0.        ],
      [1.        , 0.28169014, 0.        , ..., 0.        , 1.        ,
       0.        ],
      ...,
      [1.        , 0.53521127, 0.25      , ..., 0.        , 1.        ,
       1.        ],
      [0.5       , 0.23943662, 0.        , ..., 0.        , 1.        ,
       1.        ],
      [1.        , 0.05633803, 0.25      , ..., 0.        , 1.        ,
       1.        ]])
```